


LayerNorm bwd: triton vs persistent vs quack cute (H100, bf16) — forward + backward latency (ms, log scale; lower is better)


PyTorch Triton ours triton_persistent

forward + backward latency (ms, log)


d = 128



d = 256




d = 384



forward + backward latency (ms, log)

d = 512



d = 768

