

LayerNorm bwd: triton vs persistent vs quack cute (H100,bf16) — forward latency (ms, log scale; lower is better)

PyTorch Triton ours triton_persistent

