

LayerNorm baselines (H100, bf16) — forward + backward latency (ms, log scale; lower is better)

PyTorch Triton cuEquivariance ours

