

LayerNormLinear (H100, bf16): PyTorch / TE / Triton / ours — forward + backward latency (ms, log scale; lower is better)

