

LayerNormLinear (H100, bf16): PyTorch / Triton / Triton / ours — forward latency (ms, log scale; lower is better)

PyTorch Triton TransformerEngine ours

