



LayerNormLinear (H100, bf16): PyTorch / TE / Triton / ours — forward: speedup vs PyTorch default (higher is better)

PyTorch Triton TransformerEngine ours


d = 128




d = 256



d = 384



d = 512



d = 768

