

v6 (split back) L-scaling vs baselines (B=1, bf16, no mask; pytorch omitted)

