


Bias-only triangle attention (single-dir): triton vs torch.compile (H100, bf16) — forward + backward latency (ms, log scale; lower is better)


PyTorch Triton

forward + backward latency (ms, log)

d = 128



d = 256



d = 512

