


Bidirectional bias-only triangle attention: triton vs torch.compile (H100, bf16) — forward + backward latency (ms, log scale; lower is better)


PyTorch Triton

forward + backward latency (ms, log)

**d = 128**



**d = 256**



**d = 512**

